

INTRODUCTION:

The problem posed is prediction of continuous dependent variable i.e. interest rate from a suitable combination of independent variables, the general approach to the problem is trying to build a linear relational model as a basic benchmark model and later on building more generalized models to improve the accuracy of predictions by simultaneously aiming to reduce the error rate.

Loading the helper packages.

```
In [386]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import numpy as np
import warnings
warnings.filterwarnings('ignore')
```

Loading the data into pandas and getting a feel for it. As we can see the data needs a lot of cleaning before proceeding to analysis.

```
In [484]: cmpltddf=pd.read_csv('C:/Users/Raj/Desktop/sf work assignment.6.8.2016 (1)/Data for
```

```
In [485]: cmpltddf.head()
```

```
Out[485]:
```

	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10	...	X23
0	11.89%	54734	80364	\$25,000	\$25,000	\$19,080	36 months	B	B4	NaN	...	2/1/1
1	10.71%	55742	114426	\$7,000	\$7,000	\$673	36 months	B	B5	CNN	...	10/1/1
2	16.99%	57167	137225	\$25,000	\$25,000	\$24,725	36 months	D	D3	Web Programmer	...	6/1/2
3	13.11%	57245	138150	\$1,200	\$1,200	\$1,200	36 months	C	C2	city of beaumont texas	...	1/1/1
4	13.57%	57416	139635	\$10,800	\$10,800	\$10,692	36 months	C	C3	State Farm Insurance	...	12/1/1

5 rows × 32 columns

we can see the dimensions of our data below, we have 400000 rows and 32 columns

```
In [486]: cmpltddf.shape
```

```
Out[486]: (400000, 32)
```

Let us see how many columns have missing values

```
In [487]: cmpltddf.isnull().sum()
```

```
Out[487]: X1      61010
          X2         1
          X3         1
          X4         1
          X5         1
          X6         1
          X7         1
          X8      61270
          X9      61270
          X10     23982
          X11         1
          X12     61361
          X13     61028
          X14         1
          X15         1
          X16    276439
          X17         1
          X18        18
          X19         1
          X20         1
          X21         1
          X22         1
          X23         1
          X24         1
          X25    218802
          X26    348845
          X27         1
          X28         1
          X29         1
          X30        267
          X31         1
          X32         1
          dtype: int64
```

DATA CLEANING

There are a lot of missing values in each column, cleaning the columns before imputing missing values. We can see right away that columns X16, X25 and X26 have more than 50% of missing data we can remove those columns right away, but we will keep them for now and see if we can fill those missing values logically using the remaining data, if we cannot we will remove them later on.

```

In [488]: #cleaning data.
cmpltdf.dtypes
cmpltdf.X1=cmpltdf.X1.str.strip('%')
cmpltdf[['X4','X5','X6']]=cmpltdf[['X4','X5','X6']].apply(lambda x:x.str.strip('$'
                                                         replace(',',' ',regex=True)

cmpltdf.X7=cmpltdf.X7.str.rstrip('months')
#assigning <1 as 0.5 and 10+ years as 10,converting n/a to 0 for column X11
cmpltdf.X11=cmpltdf.X11.str.rstrip('+ years').replace('< 1','0.5',regex=True)
cmpltdf.X11=cmpltdf.X11.str.replace('n/','0')
cmpltdf.X14=cmpltdf.X14.str.rstrip('- income source')
#converting the date column into datetime.
cmpltdf.X15=pd.to_datetime(cmpltdf.X15,format='%m/%d/%Y')
cmpltdf['X23']=pd.to_datetime(cmpltdf.X23,format='%m/%d/%Y')
cmpltdf.X30=cmpltdf.X30.str.strip('%')
cmpltdf.head()

```

Out[488]:

	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10	...	X23	X24	X25
0	11.89	54734	80364	25000	25000	19080	36	B	B4	NaN	...	1994-02-01	0	NaN
1	10.71	55742	114426	7000	7000	673	36	B	B5	CNN	...	2000-10-01	0	NaN
2	16.99	57167	137225	25000	25000	24725	36	D	D3	Web Programmer	...	2000-06-01	0	41
3	13.11	57245	138150	1200	1200	1200	36	C	C2	city of beaumont texas	...	1985-01-01	0	64
4	13.57	57416	139635	10800	10800	10692	36	C	C3	State Farm Insurance	...	1996-12-01	1	58

5 rows × 32 columns



converting the dtypes of each column appropriately.

```
In [489]: cmpltddf.dtypes
```

```
Out[489]: X1          object
          X2          float64
          X3          float64
          X4          object
          X5          object
          X6          object
          X7          object
          X8          object
          X9          object
          X10         object
          X11         object
          X12         object
          X13         float64
          X14         object
          X15         datetime64[ns]
          X16         object
          X17         object
          X18         object
          X19         object
          X20         object
          X21         float64
          X22         float64
          X23         datetime64[ns]
          X24         float64
          X25         float64
          X26         float64
          X27         float64
          X28         float64
          X29         float64
          X30         object
          X31         float64
          X32         object
dtype: object
```

```
In [393]: #typecasting int numeric.
cmpltdf[['X1', 'X4','X5','X6','X13','X21','X22','X24','X25','X26','X27',
        'X28','X29','X30','X31']] = cmpltdf[['X1', 'X4','X5','X6','X13',
        'X21','X22','X24','X25','X26','X27','X
cmpltdf[['X7','X8','X9','X11','X12','X14','X17','X20','X32']] = cmpltdf[[
        'X7','X8','X9','X11','X12','X14','X17','X20','X32']].apply(lambda x:x.asty
cmpltdf.dtypes
```

```
Out[393]: X1          float64
X2          float64
X3          float64
X4          float64
X5          float64
X6          float64
X7          category
X8          category
X9          category
X10         object
X11         category
X12         category
X13         float64
X14         category
X15         datetime64[ns]
X16         object
X17         category
X18         object
X19         object
X20         category
X21         float64
X22         float64
X23         datetime64[ns]
X24         float64
X25         float64
X26         float64
X27         float64
X28         float64
X29         float64
X30         float64
X31         float64
X32         category
dtype: object
```

We can safely remove the columns X16 and X18 from our analysis as these columns contain text data in the form of comments entered by the borrowers, we will keep the X17 column i.e. the loan category as it is not random text even though it is entered by borrower. Similarly we will drop column X10 even though it look like an important column there are a lot of levels to it which cannot be categorized.

```
In [394]: cmpltdf.drop(['X16','X18','X10'],inplace = True,axis=1)
```

EXPLORATORY DATA ANALYSIS & VISUALIZATION:


In [395]:

```
cmpltdf.X1.hist()
cmpltdf.X1.describe()
#sns.boxplot(cmpltdf.X1)
#cmpltdf.loc[(cmpltdf.X1>24),['X1','X22','X4','X30','X29','X31']]
#cmpltdf.loc[(cmpltdf.X1>24),['X1','X16','X17','X18','X29','X31']]
#cmpltdf.X17.value_counts()
```

Out[395]:

count	338990.000000
mean	13.946271
std	4.377951
min	5.420000
25%	10.990000
50%	13.680000
75%	16.780000
max	26.060000


Name: X1, dtype: float64



First exploring the target variable, the target variable seems normally distributed. Next let us visualize the variable X8 which is a categorical variable and also let us see the impact it has on our target variable X1. violin plot is a good way to assess the effect of a categorical variable on other variable, the below violin plots show that the mean of the target variable is highly impacted by the grouping of levels in the target variables. The long tails in the violin plots hints outliers.

```
In [396]: #sns.countplot(x=cmltddf.X8, palette='Blues_d')
sns.violinplot(x="X8", y="X1", data=cmltddf)
```

```
Out[396]: <matplotlib.axes._subplots.AxesSubplot at 0xe6296c50>
```




Now that we know that the categorical variable X8 is going to be important in our analysis let us visualize the variable X9 which is Similar to the variable X8 provides granular level details then X8. Let us plot a facet grid plot to see the influence of variable X9 on the target variable at a much granular level.

```
In [397]: #Facet grid Plot
#g = sns.factorplot(x="X8", y="X1", hue="X9", data=cmltddf, kind="violin", size=20.

df = cmltddf.assign(X9=cmltddf.X9.astype(object)).sort("X9")
grid=sns.FacetGrid(df,col='X8', col_wrap=4,size=4)
grid.map(sns.boxplot,'X9','X1')
```

```
Out[397]: <seaborn.axisgrid.FacetGrid at 0x70126f28>
```



The variable X9 is a potential data leak as it explains almost all the variance in the target variable, this leak can be exploited to impute the missing values in the columns X1 and X9. The missing values in X1 can be imputed by the grouped means of X1 by X9 and similarly missing values in X1 can be imputed by grouping on X9. The code for imputing missing values in both the columns is below..

```
In [398]: #mean values after imputaion
cmpltddf['X9'] = cmpltddf.groupby(['X1'])['X9'].transform(lambda x: x.fillna(method=
cmpltddf['X1'].groupby(cmpltddf.X9).mean())
```


```
Out[398]: X9
A1      6.002905
A2      6.566850
A3      7.450486
A4      7.828795
A5      8.703269
B1      9.730468
B2     10.772742
B3     11.699056
B4     12.472588
B5     13.069878
C1     13.677812
C2     14.280730
C3     14.837854
C4     15.412472
C5     16.042135
D1     16.695921
D2     17.297698
D3     17.835426
D4     18.403431
D5     19.073152
E1     19.560820
E2     20.220477
E3     20.821978
E4     21.494988
E5     22.138226
F1     22.919932
F2     23.353235
F3     23.919594
F4     24.221877
F5     24.515928
G1     24.950680
G2     25.080610
G3     25.335682
G4     25.209003
G5     25.433292
Name: X1, dtype: float64
```

The Means did not change after imputation implies imputation was succesful.

similarly we will fill the missing values in X1 with the grouping of column X9.


```
In [399]: cmpltddf.X1=cmpltddf.groupby(['X9'])['X1'].transform(lambda x: x.fillna(method='ffil
cmpltddf.X1.hist()
cmpltddf.X1.isnull().sum()
```

Out[399]: 9404



```
In [400]: #cmpltddf.isnull().sum()
```

The histogram above is same as the histogram before filling the missing values, this implies that missing values imputation was successful

We can safely drop the column X8 as we have a much more granular column in X9 which can be used.

```
In [401]: cmpltddf.drop(['X8'],inplace = True,axis=1)
```

the column X12 has the most missing values. The missing values can be imputed by mode but there can be a more logical way to impute it, let us consider the effect of variable X17 on X12 with a simple groupby command, and we can see that there is influence of variable X17 on X12, let us anyway do a chi squared test to confirm that.

```
In [402]: from scipy.stats import chi2_contingency
chisrlts=chi2_contingency(pd.crosstab(cmpltddf.X12,cmpltddf.X17))
print float(chisrlts[1])
```

0.0

The p-value of Chi Square test evaluates our hypothesis that the variable X17 has influence on X12 so we can now impute the missing values in X12 with the means of X12 grouped by X17.

```
In [403]: cmpltddf.X12=cmpltddf.groupby(['X17'])['X12'].transform(lambda x: x.fillna(method='f
cmpltddf.X12.isnull().sum()
```

Out[403]: 1

Now let us fill the missing values in the column X13, the distribution looks skewed so it is better if we impute the values with median of the column.

```
In [404]: #cmltddf.isnull().sum()  
#len(cmltddf.X13.value_counts())  
#cmltddf.X13.hist().set(xlim=(0, max(cmltddf.X13)))  
cmltddf.X13.describe()
```

```
Out[404]: count      338972.000000  
mean        73160.149695  
std         55867.696483  
min         3000.000000  
25%         45000.000000  
50%         63000.000000  
75%         88200.000000  
max         750000.000000  
Name: X13, dtype: float64
```

```
In [405]: cmltddf.X13=cmltddf.X13.fillna(cmltddf.X13.median())  
#cmltddf.isnull().sum()
```

Now let us look into the column X30 which has only 200 missing values, these missing values can be filled with the mean of the column

```
In [406]: cmltddf.X30.describe()
```

```
Out[406]: count      399733.000000  
mean         56.279059  
std          23.734198  
min           0.000000  
25%          39.500000  
50%          57.800000  
75%          74.900000  
max          892.300000  
Name: X30, dtype: float64
```

```
In [407]: cmltddf.X30=cmltddf.X30.fillna(cmltddf.X30.mean())
```

Dropping X25,X26 as more than half data is missing values

```
In [408]: cmltddf.drop(['X25','X26'], inplace =True , axis=1)
```

```
In [409]: #moving into a newdf  
c1ndf=cmltddf.copy()
```

now that we are done with cleaning the data and imputing missing values we can now go ahead and remove the remaining missing values as there are no logical ways to impute the remaining missing values.

```
In [410]: cln_df.dropna(inplace=True)
```

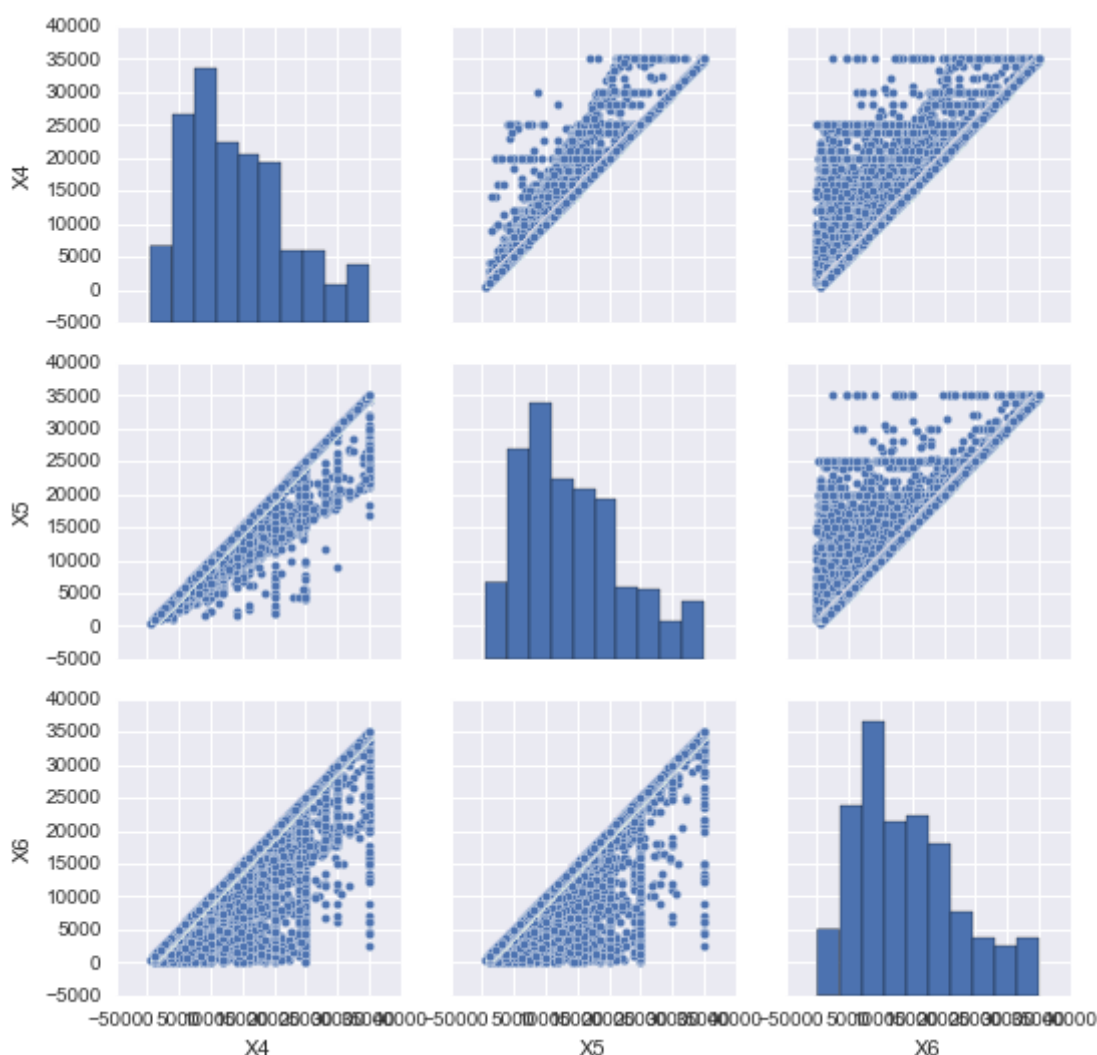
```
In [411]: #dropping X2 abd X3 as they are just ID's
cln_df.drop(['X2','X3'], inplace =True , axis=1)
```

FEATURE SELECTION:

Now we will go ahead with selecting important variables for our model.

```
In [412]: sns.pairplot(cln_df[['X4','X5','X6']])
```

```
Out[412]: <seaborn.axisgrid.PairGrid at 0x13135ec50>
```



all the three columns are linearly related with each other and we will not need all three in predicting the target variable as it may lead to multicollinearity issue.

```
In [413]: #splitting the date columns
cldf.dtypes
cldf['X15_yr']=cldf.X15.dt.year
cldf['X15_mn']=cldf.X15.dt.month
cldf['X23_yr']=cldf.X23.dt.year
cldf['X23_mn']=cldf.X23.dt.month
```

```
In [432]: #converting to category.
cldf[['X15_yr','X15_mn','X23_yr','X23_mn','X24']]=cldf[['X15_yr','X15_mn','X23_yr','X23_mn','X24']].apply(lambda x:x.astype('category'))
```

```
In [415]: #dropig the actual date columns
cldf.drop(['X15','X23'], inplace =True , axis=1)
cldf.columns
```

```
Out[415]: Index([u'X1', u'X4', u'X5', u'X6', u'X7', u'X9', u'X11', u'X12', u'X13',
u'X14', u'X17', u'X19', u'X20', u'X21', u'X22', u'X24', u'X27', u'X28',
u'X29', u'X30', u'X31', u'X32', u'X15_yr', u'X15_mn', u'X23_yr',
u'X23_mn'],
dtype='object')
```

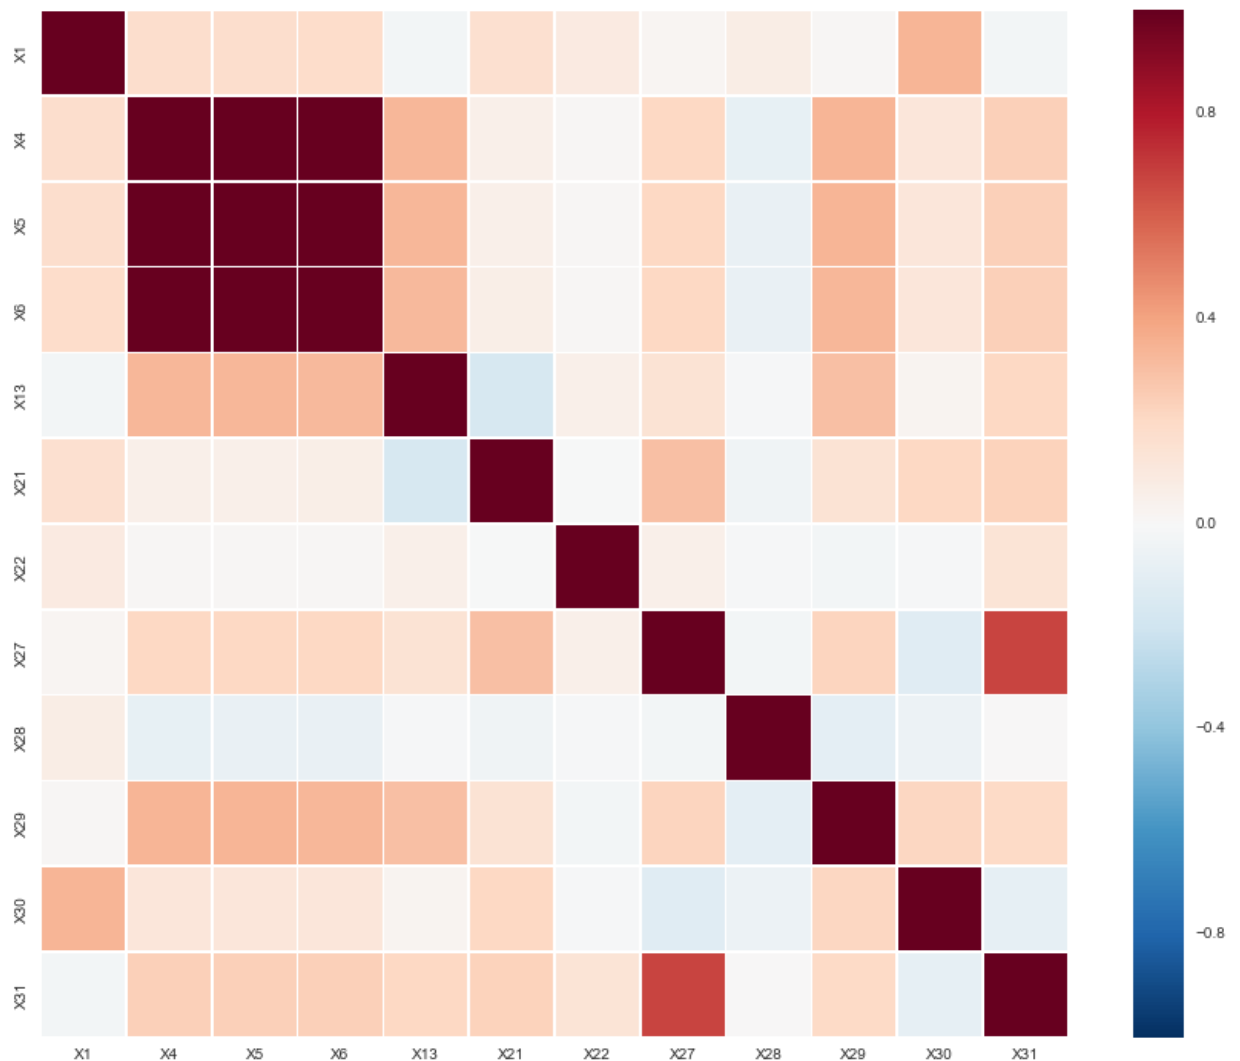
As the problem is prediction of continuous values it can be helpful for the further analysis to know which variables co-relate with the target variable and also amongst themselves. The correlations can be determined by plotting a simple heat map of all the continuous variables.

```
In [416]: corr=cldf.corr()
corr
```

```
Out[416]:
```

	X1	X4	X5	X6	X13	X21	X22	X27
X1	1.000000	0.177179	0.178194	0.180319	-0.030587	0.158253	0.091447	0.019927
X4	0.177179	1.000000	0.998346	0.994669	0.329389	0.060687	0.008901	0.204311
X5	0.178194	0.998346	1.000000	0.996645	0.328802	0.062195	0.009507	0.205256
X6	0.180319	0.994669	0.996645	1.000000	0.327191	0.065827	0.010347	0.206224
X13	-0.030587	0.329389	0.328802	0.327191	1.000000	-0.167327	0.055727	0.141342
X21	0.158253	0.060687	0.062195	0.065827	-0.167327	1.000000	-0.002156	0.304267
X22	0.091447	0.008901	0.009507	0.010347	0.055727	-0.002156	1.000000	0.062031
X27	0.019927	0.204311	0.205256	0.206224	0.141342	0.304267	0.062031	1.000000
X28	0.075079	-0.078600	-0.078006	-0.076488	-0.014786	-0.046337	-0.008924	-0.030579
X29	0.008478	0.336389	0.335996	0.334533	0.300258	0.147676	-0.030223	0.224003
X30	0.343323	0.117269	0.118605	0.120844	0.030842	0.206654	-0.011470	-0.118882
X31	-0.028160	0.237330	0.237208	0.237512	0.203970	0.229384	0.133531	0.677413

```
In [417]: f, ax = plt.subplots(figsize=(15, 12))
sns.heatmap(corr,linewidths=.5, ax=ax);
```



The columns X4, X5, X6 are all positively correlated with each other this implies we cannot use all these three or any two variables together in our models as they can lead to multicollinearity issue which is common in linear models. As we can see there are no variables which are strongly correlated with the target variable except for X30 which is slightly correlated, this gives an indication that the remaining variables can be ignored from our analysis, even though they seem not significant they can be further analyzed before being ignoring completely.

Extracting important features using multiple regression First let us extract all the columns with dtypes numeric

```
In [418]: #We will also remove the column X19 as it is not very usefull for our analysis
cldf.drop(['X19'], inplace =True , axis=1)
```

```
In [419]: cln_df.dtypes
```

```
Out[419]: X1          float64
          X4          float64
          X5          float64
          X6          float64
          X7          category
          X9          category
          X11         category
          X12         category
          X13         float64
          X14         category
          X17         category
          X20         category
          X21         float64
          X22         float64
          X24         category
          X27         float64
          X28         float64
          X29         float64
          X30         float64
          X31         float64
          X32         category
          X15_yr       category
          X15_mn       category
          X23_yr       category
          X23_mn       category
          dtype: object
```

```
In [420]: #first let us split our variable into categories and numericals so that we can wor
          catgrydf=cln_df.loc[:,cln_df.dtypes.isin(['category'])]
          quantdf=cln_df.loc[:,cln_df.columns-catgrydf.columns]
          catgrydf['X1']=cln_df.X1
```

In [421]: *# we can first do a simple groupby function to find out if the means vary by each #for variable X9.*

```
#catgrydf.columns
for i in range(len(catgrydf.columns)):
    print catgrydf.groupby(catgrydf.columns[i])['X1'].mean()
#the group means are changing
```

```
X7
36    12.736770
60    17.225859
Name: X1, dtype: float64
X9
A1     6.001833
A2     6.564321
A3     7.446712
A4     7.829278
A5     8.703836
B1     9.730990
B2    10.772892
B3    11.698913
B4    12.474209
B5    13.069579
C1    13.680796
C2    14.282746
C3    14.838568
C4    15.415817
C5    16.044457
```

The group means are changing only for a the categorical columns X9 and X7, for the rest there are closely same. which implies only X7 and X9 are important variables.

Label encoding the variables so that it can be used for further analysis.

In [422]: *#LabelEncoding('X17',catgrydf)*
#catgrydf.loc[['X9','X12','X14','X17','X20','X32']].apply(Lambda x:LabelEncoding(x

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
catgrydf['X9_enc']=pd.Series(le.fit_transform(cmpltdf.X9))
catgrydf['X12_enc']=pd.Series(le.fit_transform(cmpltdf.X12))
catgrydf['X14_enc']=pd.Series(le.fit_transform(cmpltdf.X14))
catgrydf['X17_enc']=pd.Series(le.fit_transform(cmpltdf.X17))
catgrydf['X20_enc']=pd.Series(le.fit_transform(cmpltdf.X20))
catgrydf['X32_enc']=pd.Series(le.fit_transform(cmpltdf.X32))
```

In [423]: *#Converting the new types into catagorey*
catgrydf[['X9_enc','X12_enc','X14_enc','X17_enc','X20_enc','X32_enc']]=catgrydf[['
'X9_enc','X12_enc','X14_enc','X17_enc','X20_enc','X32_enc']].apply(lambda

```
In [424]: #deleting all the converted variables.
catgrydf.drop(['X9','X12','X14','X17','X20','X32'],axis=1,inplace=True)
catgrydf.drop(['X1'],axis=1,inplace=True)
```

```
In [425]: #merging and copying into a new dataframe
newdf=pd.concat([catgrydf,quantdf],axis=1)
newdf.isnull().sum()
newdf.shape
```

```
Out[425]: (390539, 25)
```

Splitting the data into dependent and independent variables

```
In [426]: train=newdf
Y_train_data=train.X1
X_train_data=train.drop(['X1'],1)
X_train_data.columns
```

```
Out[426]: Index([u'X7', u'X11', u'X24', u'X15_yr', u'X15_mn', u'X23_yr', u'X23_mn',
u'X9_enc', u'X12_enc', u'X14_enc', u'X17_enc', u'X20_enc', u'X32_enc',
u'X13', u'X21', u'X22', u'X27', u'X28', u'X29', u'X30', u'X31', u'X4',
u'X5', u'X6'],
dtype='object')
```

```
In [427]: #train test split
from sklearn.preprocessing import StandardScaler
from sklearn.cross_validation import train_test_split
scaler=preprocessing.StandardScaler()
X_train, X_test, Y_train, Y_test = train_test_split(X_train_data, Y_train_data, te
#scaler.fit_transform(X_train_data)
```

```
In [428]: #scaler.fit(X_train_data)
X_train_data.head()
```

```
Out[428]:
```

	X7	X11	X24	X15_yr	X15_mn	X23_yr	X23_mn	X9_enc	X12_enc	X14_enc	...	X21	X
0	36	0.5	0	2016	8	1994	2	9	6	1	...	19.48	0
1	36	0.5	0	2016	5	2000	10	10	6	2	...	14.29	0
2	36	1	0	2016	8	2000	6	18	6	1	...	10.50	0
3	36	10	0	2016	3	1985	1	12	5	2	...	5.47	0
4	36	6	1	2016	11	1996	12	13	6	2	...	11.63	0

5 rows × 24 columns



The above inferences made about the important features can be tested with various statistical tests and other methods. First thing that can be done is compute a simple Pearson correlations amongst all the variable and get the most important variables, The f_regression function from the

preprocessing package in sklearn can be very efficient in computing the p-values but first we need to get the data ready for analysis.

In [429]:

```
from sklearn.feature_selection import f_regression
from sklearn.feature_selection import SelectKBest
from sklearn.preprocessing import StandardScaler
scaler=preprocessing.StandardScaler()
featureSelector = SelectKBest(score_func=f_regression,k=5)
featureSelector.fit(scaler.fit_transform(X_train_data),scaler.fit_transform(Y_train_data))
X_train_data.loc[:,featureSelector.get_support()].columns, featureSelector.get_support()
```

Out[429]: (Index([u'X7', u'X24', u'X9_enc', u'X14_enc', u'X30'], dtype='object'),
array([True, False, True, False, False, False, False, True, False,
 True, False, False, False, False, False, False, False, False,
 False, True, False, False, False, False], dtype=bool))

The results of a simple correlation test has given the five most important variables as specified in the code, of this five using only the two most important variables to build a base model, If using more than two variables multicollinearity issue creeps into the model

MODEL BUILDING:

Using statsmodels package in python to build the base model as it provides various statistics which can be interactive in feature selection.

In [474]:

```
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
all_columns = "+".join(X_train[['X7','X9_enc']])
my_formula = "Y_train~" + all_columns
lm=smf.ols(formula=my_formula, data=X_train).fit()
#lm = smf.ols(np.array(Y_train_data),np.array())
```

In [475]: `print lm.summary()`

OLS Regression Results

```

=====
Dep. Variable:          Y_train    R-squared:                0.965
Model:                  OLS        Adj. R-squared:           0.965
Method:                 Least Squares    F-statistic:             2.141e+05
Date:                  Tue, 05 Jul 2016    Prob (F-statistic):       0.00
Time:                  11:39:04    Log-Likelihood:          -3.3395e+05
No. Observations:      273377    AIC:                     6.680e+05
Df Residuals:          273341    BIC:                     6.683e+05
Df Model:              35
Covariance Type:       nonrobust
=====

```

```

=====
=
              coef      std err          t      P>|t|      [95.0% Conf. In
t.]
-----
-
Intercept      6.0000      0.010     572.737      0.000      5.979      6.02
0
X7[T. 60 ]    -0.1124      0.004    -27.707      0.000     -0.120     -0.10
4
X9_enc[T.2]     0.5642      0.015     38.175      0.000      0.535      0.59
3
X9_enc[T.3]     1.4469      0.014    101.614      0.000      1.419      1.47
5
X9_enc[T.4]     1.8307      0.013    140.267      0.000      1.805      1.85
6
X9_enc[T.5]     2.7134      0.013    212.094      0.000      2.688      2.73
8
X9_enc[T.6]     3.7352      0.013    295.934      0.000      3.710      3.76
0
X9_enc[T.7]     4.7815      0.012    387.982      0.000      4.757      4.80
6
X9_enc[T.8]     5.7080      0.012    474.561      0.000      5.684      5.73
2
X9_enc[T.9]     6.4897      0.012    535.046      0.000      6.466      6.51
3
X9_enc[T.10]    7.0812      0.012    567.158      0.000      7.057      7.10
6
X9_enc[T.11]    7.7076      0.012    622.264      0.000      7.683      7.73
2
X9_enc[T.12]    8.3176      0.012    669.427      0.000      8.293      8.34
2
X9_enc[T.13]    8.8814      0.013    707.982      0.000      8.857      8.90
6
X9_enc[T.14]    9.4612      0.013    745.664      0.000      9.436      9.48
6
X9_enc[T.15]   10.1027      0.013    788.637      0.000     10.078     10.12
8
X9_enc[T.16]   10.7490      0.013    818.401      0.000     10.723     10.77
5
X9_enc[T.17]   11.3372      0.013    847.890      0.000     11.311     11.36
3
X9_enc[T.18]   11.8786      0.014    863.904      0.000     11.852     11.90
6
X9_enc[T.19]   12.4575      0.014    888.149      0.000     12.430     12.48
5

```

X9_enc[T.20] 9	13.1298	0.015	896.023	0.000	13.101	13.15
X9_enc[T.21] 5	13.6340	0.016	866.855	0.000	13.603	13.66
X9_enc[T.22] 5	14.2837	0.016	900.149	0.000	14.253	14.31
X9_enc[T.23] 4	14.8709	0.017	880.495	0.000	14.838	14.90
X9_enc[T.24] 3	15.5989	0.018	885.317	0.000	15.564	15.63
X9_enc[T.25] 7	16.2302	0.019	874.883	0.000	16.194	16.26
X9_enc[T.26] 6	17.0268	0.020	850.956	0.000	16.988	17.06
X9_enc[T.27] 8	17.4839	0.022	779.896	0.000	17.440	17.52
X9_enc[T.28] 7	17.9924	0.023	782.605	0.000	17.947	18.03
X9_enc[T.29] 0	18.2796	0.026	710.419	0.000	18.229	18.33
X9_enc[T.30] 7	18.6392	0.030	629.840	0.000	18.581	18.69
X9_enc[T.31] 6	19.0286	0.034	554.567	0.000	18.961	19.09
X9_enc[T.32] 8	19.1926	0.039	498.133	0.000	19.117	19.26
X9_enc[T.33] 2	19.4044	0.045	435.395	0.000	19.317	19.49
X9_enc[T.34] 3	19.2168	0.054	355.574	0.000	19.111	19.32
X9_enc[T.35] 7	19.5301	0.060	326.407	0.000	19.413	19.64

```
=====
Omnibus:                124432.943    Durbin-Watson:                1.995
Prob(Omnibus):           0.000    Jarque-Bera (JB):            1707177.122
Skew:                    -1.833    Prob(JB):                     0.00
Kurtosis:                14.680    Cond. No.                     46.7
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [442]: lm.predict(X_test[['X7','X9_enc']])
```

```
Out[442]: array([ 8.7133385 ,  8.7133385 , 19.63401416, ..., 16.10268529,
                  9.73513969, 12.37729437])
```

The summary shows that the R-squared value and adjusted R-squared values are high suggesting goodness of the fit. Also there is no trace of multicollinearity issue as none of the coefficients are small or negligible and also the summary suggest that all the tests are passed without any warnings. Now assessing the error on the train data.

```
In [476]: #lm.predict(X_train[[u'X7', u'X9_enc', u'X14_enc', u'X24']])
          np.sqrt(lm.mse_resid)
```

```
Out[476]: 0.82094297938615446
```

The train error is 0.821 on train data similar error can also be calculated on the test data. Writing a simple function called metric which will compute the error metric on test

```
In [299]: def metric(dfXtest,dfYtest,clf,standardize=False):
          if standardize:
              from sklearn.preprocessing import StandardScaler
              scaler=preprocessing.StandardScaler()
              dfXtest=scaler.fit_transform(dfXtest)
          else:
              dfXtest
          from sklearn.metrics import mean_squared_error
          predicted=clf.predict(dfXtest)
          return np.sqrt(mean_squared_error(dfYtest,predicted)),predicted,dfYtest
```

```
In [477]: testerror,predictedvals,actual=metric(X_test[ ['X7','X9_enc']],Y_test,lm)
          testerror
```

```
Out[477]: 0.82300565247332569
```

The error on the test is 0.815, the basic linear classifier has slightly missed in accuracy but it has done fairly well being the simplest model. A more regularized model might be a good fit because if one notices the magnitude of the coefficients of our simple model they are huge which might be the reason for the error not being accurate and the magnitude of coefficients can be reduced by adding penalty which the regularization models like ridge and lasso can take care of. We shall see if fitting a regularized model can solve the problem of accuracy as well.

In order to implement regularized model scikit learn is more suitable as it allows doing cross validation to find the best parameters the below function CV is a simple implementation of sklearn GridSearchCV module which finds the best parameters after going over each fold.

```
In [456]: def cv(clf, parameters, Xdf, ydf,n_folds=5):
          from sklearn.grid_search import GridSearchCV
          gs = GridSearchCV(clf, param_grid=parameters, cv=n_folds)
          gs.fit(X_train, Y_train)
          print "BEST PARAMS ARE", gs.best_params_
          #print "Best Score", gs.grid_scores_
          best = gs.best_estimator_
          #cv_score=gs.grid_scores_
          return best
```

The below function regression is the main function which uses the earlier described function metric and cv to fit the best parameters found in cv over the whole training data and give the error metric on train.

```
In [276]: def regression(clf, parameters, X, Y,n_folds=5,train_size=0.7,features=False,stand
          if features:
              subdf=X[args]
          else:
              subdf=X
          if standardize:
              from sklearn.preprocessing import StandardScaler
              scaler=preprocessing.StandardScaler()
              subdfstd=scaler.fit_transform(subdf)
          else:
              subdfstd=subdf
          Xtr=subdfstd
          Ytr=Y
          X_train, X_test, Y_train, Y_test = train_test_split(Xtr, Ytr, train_size=train
          clf = cv(clf, parameters, X_train, Y_train,n_folds=5)
          clf=clf.fit(X_train, Y_train)
          trainRMSE,xjunk,yjunk = metric(X_train, Y_train,clf,standardize=False)
          print "Rmse on Train: %0.5f" % (trainRMSE)
          return clf, X_train, Y_train, X_test, Y_test,trainRMSE
```

The below function dum, creates dummies(one hot encoding) suitable for sklearn implementations

```
In [279]: def dum(df):
          dum=pd.get_dummies(df)
          return dum.iloc[:,1:]
```

```
In [443]: #preparing data for regularization
          X=pd.concat([dum(X_train_data['X7']),dum(X_train_data['X9_enc'])],axis=1)
```

Importing the ridge regression module and tuning parameters using the regression function. The ridge regression applies L1 penalty on the coefficients with large magnitudes. The best parameter of alpha(L1) is 100, on evaluating the error on test data we can see the accuracy has increase when compared to the benchmark model.

```
In [380]: from sklearn.linear_model import Ridge
          clfRIDGE=Ridge()
          #clfRIDGE.get_params()
```

```
In [478]: RDGmdl,xtRdg,ytRdg,xteRdg,yteRdg,trainerrorRdg=regression(clfRIDGE,{'alpha':[0.1,1
                                                                    train_size=0.7,
                                                                    standardize=True)
```

```
BEST PARAMS ARE {'alpha': 100}
Rmse on Train: 0.82085
```

```
In [462]: testerrorRdg,predictedvalsRdG,actualRdg=metric(xteRdg,yteRdg,RDGmdl)
          testerrorRdg
```

```
Out[462]: 0.82701673013874522
```

```
In [446]: RDGmdl.score(xtRdg,ytRdg)
```

```
Out[446]: 0.96486867514213959
```

The R-squared value of the model is 0.96, so we can conclude that the Ridge model has the same error rate as the linear model. Moving on to tree based models, as tree based models like random forests and XGboosts work well when the dependent variables are categorical.

Building a random forest classifier, even though random forest is based on bagging technique, it is highly recommended to cross validate. Setting up the appropriate parameters and parsing the classifier to the regression function defined above.

```
In [457]: from sklearn.ensemble import RandomForestRegressor
clfRF=RandomForestRegressor()
clfRF.get_params()
```

```
Out[457]: {'bootstrap': True,
'criterion': 'mse',
'max_depth': None,
'max_features': 'auto',
'max_leaf_nodes': None,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 10,
'n_jobs': 1,
'oob_score': False,
'random_state': None,
'verbose': 0,
'warm_start': False}
```

```
In [458]: RFmdl,xtRF,ytRF,xteRF,yteRF,trainerrorRF=regression(clfRF,{'n_estimators':[200,300
'random_state': [10]]
,X,Y_train_data,n_folds=5,train_size=0.7,
standardize=False)
```

```
BEST PARAMS ARE {'max_features': 'sqrt', 'n_estimators': 300, 'n_jobs': 3, 'ran
dom_state': 10}
Rmse on Train: 0.81846
```

```
In [459]: testerrorRF,predictedvalsRF,actualRF=metric(xteRF,yteRF,RFmdl)
testerrorRF
```

```
Out[459]: 0.81935346148604105
```

```
In [460]: RFmdl.score(xtRF,ytRF)
```

```
Out[460]: 0.96494142755709289
```

As we can see the error rate has decreased significantly with the forests as expected and also the score suggests the model is a good fit the Random forests worked better than generalization methods because the most important variables are all categories

Cleaning the test data so that the trained classifiers can be used to predict the values on holdout data.

```
In [311]: test=pd.read_csv('C:/Users/Raj/Desktop/sf work assignment.6.8.2016 (1)/Holdout for
```

```
In [312]: def test_clean(df):
    #df['X1']=df['X1'].str.strip('%')
    df[['X4','X5','X6']]=df[['X4','X5','X6']].apply(lambda x:x.str.strip('$').repl
    df.X7=df.X7.str.rstrip('months')
    #assigning <1 as 0.5 and 10+ years as 10,converting n/a to 0 for column X11
    df.X11=df.X11.str.rstrip('+ years').replace('< 1','0.5',regex=True)
    df.X11=df.X11.str.replace('n/','0')
    df.X14=df.X14.str.rstrip('- income source')
    #converting the date column into datetime.
    df.X15=pd.to_datetime(df.X15,format='%m/%d/%Y')
    df['X23']=pd.to_datetime(df.X23,format='%m/%d/%Y')
    df.X30=df.X30.str.strip('%')
    df['X15_yr']=df.X15.dt.year
    df['X15_mn']=df.X15.dt.month
    df['X23_yr']=df.X23.dt.year
    df['X23_mn']=df.X23.dt.month
    #changing dtypes
    df[['X4','X5','X6','X13','X21','X22','X25','X26','X27',
        'X28','X29','X30','X31']]=df[['X4','X5','X6','X13',
        'X21','X22','X25','X26','X27','X28','X
    df[['X7','X8','X9','X11','X12','X14','X17','X20','X32','X15_yr','X15_mn','X23_
        'X7','X8','X9','X11','X12','X14','X17','X20','X32','X15_yr','X15_mn','X23_
    return df
```

```
In [313]: test=test_clean(test)
```

```
In [467]: #preparing test data
Xte=pd.concat([dum(test['X7']),dum(test['X9'])],axis=1)
```

Predicting with ridge classifier

```
In [481]: RDGpredictions=RDGmdl.predict(Xte)
RDGpredictions=pd.Series(RDGpredictions)
RDGpredictions.describe()
```

```
Out[481]: count      80000.000000
mean         15.416918
std           0.636854
min          13.944138
25%          15.037385
50%          15.651920
75%          15.946009
max          16.073439
dtype: float64
```

Predicting with Random Forest classifier.


```
In [480]: RFpredictions=RFmdl.predict(Xte)
RFpredictions=pd.Series(rftemp)
RFpredictions.describe()
```

```
Out[480]: count      80000.000000
mean         13.945600
std           4.270745
min           6.003808
25%          10.784590
50%          13.711056
75%          16.777297
max          25.650080
dtype: float64
```

```
In [482]: out=pd.concat([RDGpredictions,RFpredictions],axis=1)
out.to_csv('out.csv')
```

